


Classification And Regression Trees : A Practical Guide for Describing a Dataset

Leo Pekelis


February 2nd, 2013, Bicoastal Datafest, Stanford University

What is a Tree?

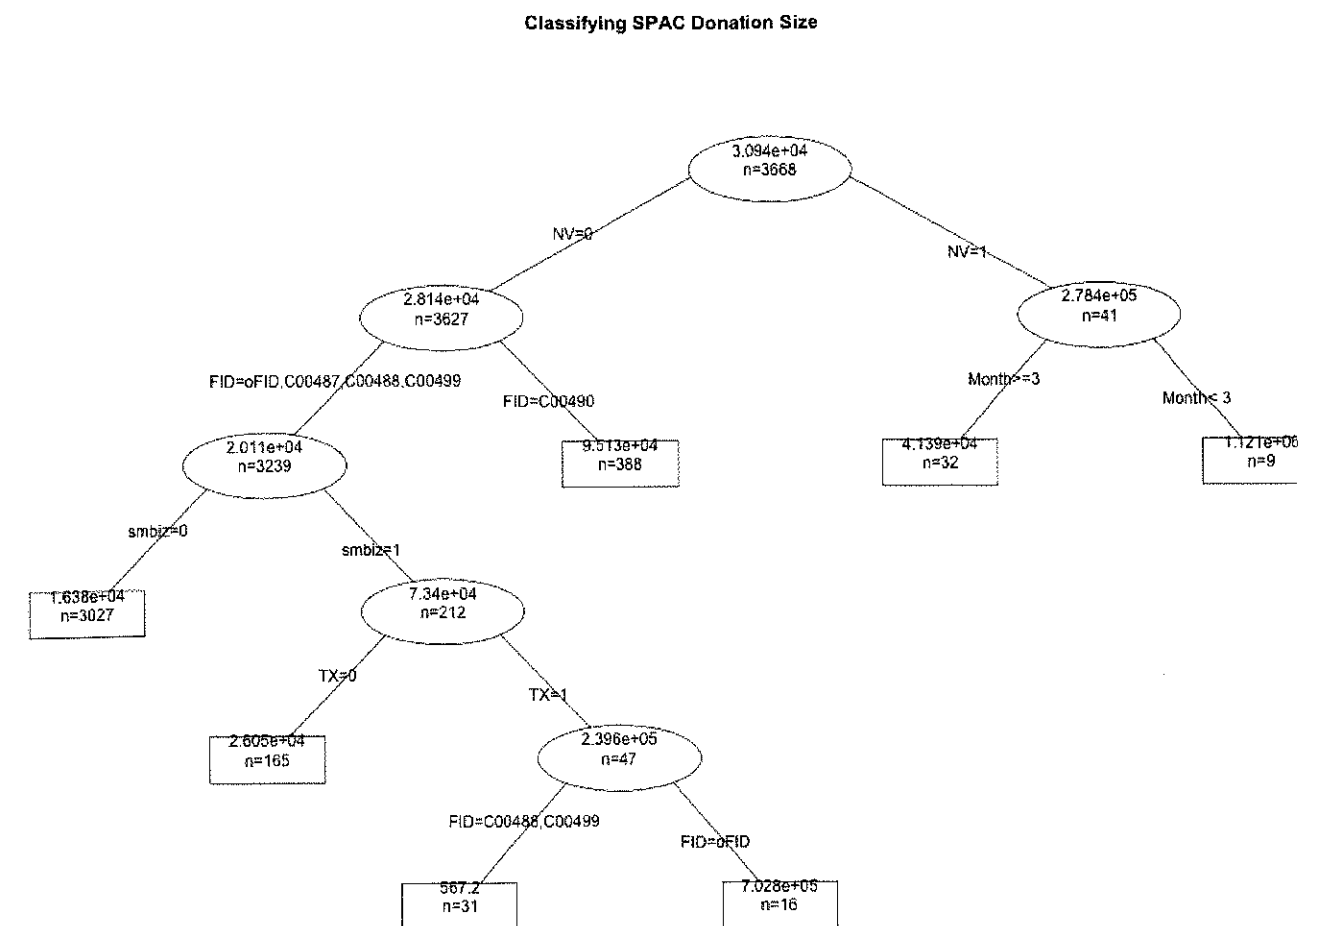


... ?

What is a (binary) Decision Tree?



What is a (binary) Decision Tree? Example



The data is all donations to SPACs in excess of \$200, by early 2012, from fec.gov

The Structural Model

- $F(x) = \sum_{i=1}^M c_m I(x \in R_m)$
- $\{R_m\}_1^M$ are subregions of the input variable space, and x is a vector of input variables.
 - Examples: $\{x_9 < 15.2\}$, $\{9 \leq x_{300} < 786 \ \& \ color = red\}$
- c_m are the estimated values of the outcome (y) in region R_m
- CART tries to minimize
 - $e(T) = \sum_{i=1}^N \left[y_i - \sum_{m=1}^M c_m I(x \in R_m) \right]^2$
 - with respect to c_m and R_m

Some Important Facts about CART

The R_m regions are disjoint and rectangular

- giving a piecewise constant approximation to the true $F(x)$

CART doesn't find the "best" regions exactly

- uses recursive partitioning, or a greedy stepwise descent

Both simplifications are to simplify a combinatorially hard problem and make it solvable in reasonable time.

- also allows for natural representations of regions as a binary decision tree

How do we run it?

```
# install the package to R
install.packages("rpart", repos = "http://cran.us.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/0m/xzr0fktj78sgl36y77z34djr0000gn/T//RtmpPULWHm/downloaded_packages
```

```
# load the library
library(rpart)

# load the dataset
load("spac.Rdata")

spac.tree = rpart(Donation ~ ., data = spac.data, cp = 10^(-6))

#### the function arguments:

# 1) formula, of the form: outcome ~ predictors

# note: outcome ~ . is 'use all other variables in data'

# 2) data: a data.frame object, or any matrix which has variables as
# columns and observations as rows

# 3) cp: used to choose depth of the tree, we'll manually prune the tree
# later and hence set the threshold very low (more on this later)

# The commands, print() and summary() will be useful to look at the tree.
# But first, lets see how big the created tree was

# The object spac.tree is a list with a number of entires that can be
# accessed via the $ symbol. A list is like a hash table.

# To see the entries in a list, use names()
names(spac.tree)
```

```
## [1] "frame"      "where"      "call"
## [4] "terms"      "cptable"    "method"
## [7] "parms"      "control"    "functions"
## [10] "numresp"    "splits"     "csplit"
## [13] "variable.importance" "y"          "ordered"
```

```
# Within spac.tree the cptable will tell us a little about the size of the
# tree
spac.tree$cptable[1:10, ]
```

```
##          CP nsplit rel error xerror  xstd
## 1  0.037317      0  1.0000  1.000  0.3477
## 2  0.016462      2  0.9254  1.078  0.3493
## 3  0.003617      6  0.8595  1.068  0.3300
## 4  0.002751      8  0.8523  1.051  0.3171
## 5  0.001581      9  0.8495  1.050  0.3170
## 6  0.001516     17  0.8369  1.064  0.3170
## 7  0.001470     21  0.8305  1.064  0.3170
## 8  0.001454     27  0.8217  1.066  0.3170
## 9  0.001432     29  0.8188  1.066  0.3170
## 10 0.001020     32  0.8145  1.069  0.3170
```

```
# ...
```

```
spac.tree$cptable[dim(spac.tree$cptable)[1] - 9:0, ]
```

```
##          CP nsplit rel error xerror  xstd
## 84 1.901e-06     169  0.7951  1.067  0.3133
```

```
# 85 1.725e-06     170  0.7951  1.067  0.3133
## 86 1.584e-06     171  0.7951  1.067  0.3133
## 87 1.188e-06     172  0.7951  1.067  0.3133
## 88 1.177e-06     173  0.7951  1.067  0.3133
## 89 1.156e-06     174  0.7951  1.067  0.3133
## 90 1.135e-06     175  0.7951  1.067  0.3133
## 91 1.129e-06     177  0.7951  1.067  0.3133
## 92 1.061e-06     179  0.7951  1.067  0.3133
## 93 1.000e-06     181  0.7951  1.067  0.3133
```

that's a lot of splits! I'm going to prune the tree to 9 splits

```
p9 = which(spac.tree$cptable[, 2] == 9)

pac.tree9 = prune(spac.tree, spac.tree$cptable[cp9, 1])

now lets look at the tree with print() and summary()

rint(spac.tree9)
```

```
# n= 3668
#
# node), split, n, deviance, yval
# * denotes terminal node
#
# 1) root 3668 1.438e+14  30940.0
# 2) NV=0 3627 9.400e+13  28140.0
# 4) FID=otherFID,C00487470,C00488403,C00499335 3239 8.088e+13  20110.0
# 8) smbiz=0 3027 2.897e+13  16380.0
# 16) blank=0 2467 1.580e+13  10930.0 *
# 17) blank=1 560 1.278e+13  40370.0 *
# 9) smbiz=1 212 5.126e+13  73400.0
# 18) TX=0 165 1.867e+12  26050.0 *
# 19) TX=1 47 4.772e+13  239600.0
# 38) FID=C00488403,C00499335 31 5.142e+06  567.2 *
# 39) FID=otherFID 16 4.252e+13  702800.0 *
# 5) FID=C00490045 388 1.117e+13  95130.0
# 10) NY=0 345 6.533e+12  82900.0 *
# 11) NY=1 43 4.176e+12  193300.0
# 22) Day< 27.5 35 2.033e+12  138000.0 *
# 23) Day>=27.5 8 1.568e+12  435000.0 *
# 3) NV=1 41 4.723e+13  278400.0
# 6) Month>=3 32 3.476e+11  41390.0 *
# 7) Month< 3 9 3.869e+13  1121000.0 *
```

```
ummary(spac.tree9)
```

```
# Call:
# rpart(formula = Donation ~ ., data = spac.data, cp = 10^(-6))
# n= 3668
```

```
##          CP nsplit rel error xerror  xstd
## 1  0.037317      0  1.0000  1.000  0.3477
## 2  0.016462      2  0.9254  1.078  0.3493
## 3  0.003617      6  0.8595  1.068  0.3300
## 4  0.002751      8  0.8523  1.051  0.3171
## 5  0.001581      9  0.8495  1.050  0.3170
```

```
# Variable importance
# Month      FID      NV      TX      tech      oil      doctor writing      smbiz
# 35         28       9       6       3       3       3       3       2
# Day        NY      blank      biz
# 2          2       1       1
```

```
# Node number 1: 3668 observations, complexity param=0.03732
# mean=3.094e+04, MSE=3.919e+10
# left son=2 (3627 obs) right son=3 (41 obs)
```

```
# Primary splits:
# NV      splits as LR,      improve=0.017660, (0 missing)
# FID     splits as LRLLL,   improve=0.012390, (0 missing)
# Month   < 5.5 to the right, improve=0.005567, (0 missing)
```

```

##      smbiz splits as LR,      improve=0.004716, (0 missing)
##      retired splits as RL,    improve=0.003653, (0 missing)
##
## Node number 2: 3627 observations,      complexity param=0.01646
## mean=2.814e+04, MSE=2.592e+10
## left son=4 (3239 obs) right son=5 (388 obs)
## Primary splits:
##      FID splits as LRLLL,    improve=0.020740, (0 missing)
##      smbiz splits as LR,      improve=0.008136, (0 missing)
##      money splits as LR,      improve=0.004718, (0 missing)
##      retired splits as RL,    improve=0.004439, (0 missing)
##      Month < 6.5 to the right, improve=0.004148, (0 missing)
## Surrogate splits:
##      UT splits as LR, agree=0.897, adj=0.036, (0 split)
##      leisure splits as LR, agree=0.893, adj=0.003, (0 split)
##
## Node number 3: 41 observations,      complexity param=0.03732
## mean=2.784e+05, MSE=1.152e+12
## left son=6 (32 obs) right son=7 (9 obs)
## Primary splits:
##      Month < 3 to the right, improve=0.17340, (0 missing)
##      Day < 7.5 to the right, improve=0.02769, (0 missing)
##      manage splits as LR,      improve=0.02717, (0 missing)
##      FID splits as RL--L,    improve=0.02251, (0 missing)
##      professional splits as RL, improve=0.01382, (0 missing)
## Surrogate splits:
##      doctor splits as LR, agree=0.805, adj=0.111, (0 split)
##      tech splits as LR, agree=0.805, adj=0.111, (0 split)
##      oil splits as LR, agree=0.805, adj=0.111, (0 split)
##      writing splits as LR, agree=0.805, adj=0.111, (0 split)
##
## Node number 4: 3239 observations,      complexity param=0.01646
## mean=2.011e+04, MSE=2.497e+10
## left son=8 (3027 obs) right son=9 (212 obs)
## Primary splits:
##      smbiz splits as LR,      improve=0.007964, (0 missing)
##      FID splits as R-LLL,    improve=0.005066, (0 missing)
##      blank splits as LR,      improve=0.003437, (0 missing)
##      TX splits as LR,        improve=0.002374, (0 missing)
##      retired splits as RL,    improve=0.002351, (0 missing)
##
## Node number 5: 388 observations,      complexity param=0.003617
## mean=9.513e+04, MSE=2.88e+10
## left son=10 (345 obs) right son=11 (43 obs)
## Primary splits:
##      NY splits as LR,        improve=0.041680, (0 missing)
##      Day < 27.5 to the left, improve=0.028980, (0 missing)
##      CA splits as RL,        improve=0.023540, (0 missing)
##      retired splits as RL,    improve=0.011260, (0 missing)
##      Month < 1.5 to the left, improve=0.007873, (0 missing)
## Surrogate splits:
##      community splits as LR, agree=0.892, adj=0.023, (0 split)
##
## Node number 6: 32 observations
## mean=4.139e+04, MSE=1.086e+10
##
## Node number 7: 9 observations
## mean=1.121e+06, MSE=4.299e+12
##
## Node number 8: 3027 observations,      complexity param=0.002751
## mean=1.638e+04, MSE=9.572e+09
## left son=16 (2467 obs) right son=17 (560 obs)
## Primary splits:
##      blank splits as LR,      improve=0.013650, (0 missing)
##      FID splits as R-LLL,    improve=0.007902, (0 missing)
##      DC splits as LR,        improve=0.007561, (0 missing)
##      retired splits as RL,    improve=0.003920, (0 missing)
##      Day < 14.5 to the left, improve=0.002814, (0 missing)
## Surrogate splits:
##      DC splits as LR, agree=0.870, adj=0.300, (0 split)
##      ZZ splits as LR, agree=0.815, adj=0.002, (0 split)
##
## Node number 9: 212 observations,      complexity param=0.01646
## mean=7.34e+04, MSE=2.418e+11

```

```

left son=18 (165 obs) right son=19 (47 obs)
Primary splits:
      TX splits as LR,      improve=0.032550, (0 missing)
      Month < 1.5 to the right, improve=0.017010, (0 missing)
      FID splits as R-LLL,    improve=0.009249, (0 missing)
      Day < 28.5 to the left, improve=0.007682, (0 missing)
      professional splits as RL, improve=0.002284, (0 missing)
Surrogate splits:
      FID splits as L-LRL, agree=0.892, adj=0.511, (0 split)
      teach splits as LR,      agree=0.783, adj=0.021, (0 split)
      oil splits as LR,        agree=0.783, adj=0.021, (0 split)

Node number 10: 345 observations
mean=8.29e+04, MSE=1.894e+10

Node number 11: 43 observations,      complexity param=0.003617
mean=1.933e+05, MSE=9.711e+10
left son=22 (35 obs) right son=23 (8 obs)
Primary splits:
      Day < 27.5 to the left, improve=0.137500, (0 missing)
      Month < 5 to the right, improve=0.062300, (0 missing)
      money splits as LR,      improve=0.012980, (0 missing)
      professional splits as RL, improve=0.010520, (0 missing)
      manage splits as LR,      improve=0.009981, (0 missing)
Surrogate splits:
      tech splits as LR, agree=0.837, adj=0.125, (0 split)

Node number 16: 2467 observations
mean=1.093e+04, MSE=6.405e+09

Node number 17: 560 observations
mean=4.037e+04, MSE=2.282e+10

Node number 18: 165 observations
mean=2.605e+04, MSE=1.131e+10

Node number 19: 47 observations,      complexity param=0.01646
mean=2.396e+05, MSE=1.015e+12
left son=38 (31 obs) right son=39 (16 obs)
Primary splits:
      FID splits as R--LL,    improve=0.109000, (0 missing)
      Day < 28.5 to the left, improve=0.043090, (0 missing)
      Month < 5 to the right, improve=0.038900, (0 missing)
      manage splits as RL,    improve=0.005604, (0 missing)
Surrogate splits:
      Month < 3.5 to the right, agree=0.787, adj=0.375, (0 split)
      biz splits as LR,        agree=0.681, adj=0.063, (0 split)

Node number 22: 35 observations
mean=1.38e+05, MSE=5.809e+10

Node number 23: 8 observations
mean=4.35e+05, MSE=1.961e+11

Node number 38: 31 observations
mean=567.2, MSE=1.659e+05

Node number 39: 16 observations
mean=7.028e+05, MSE=2.658e+12

```

```

finally, lets get a graphical representation of the tree, and save to a
png file
fig("spactree9.png", width = 1200, height = 800)
plot(spac.tree9, file = "", title = "Classifying SPAC Donation Size, 9 splits",
     bp = 18)
sv.off()

```

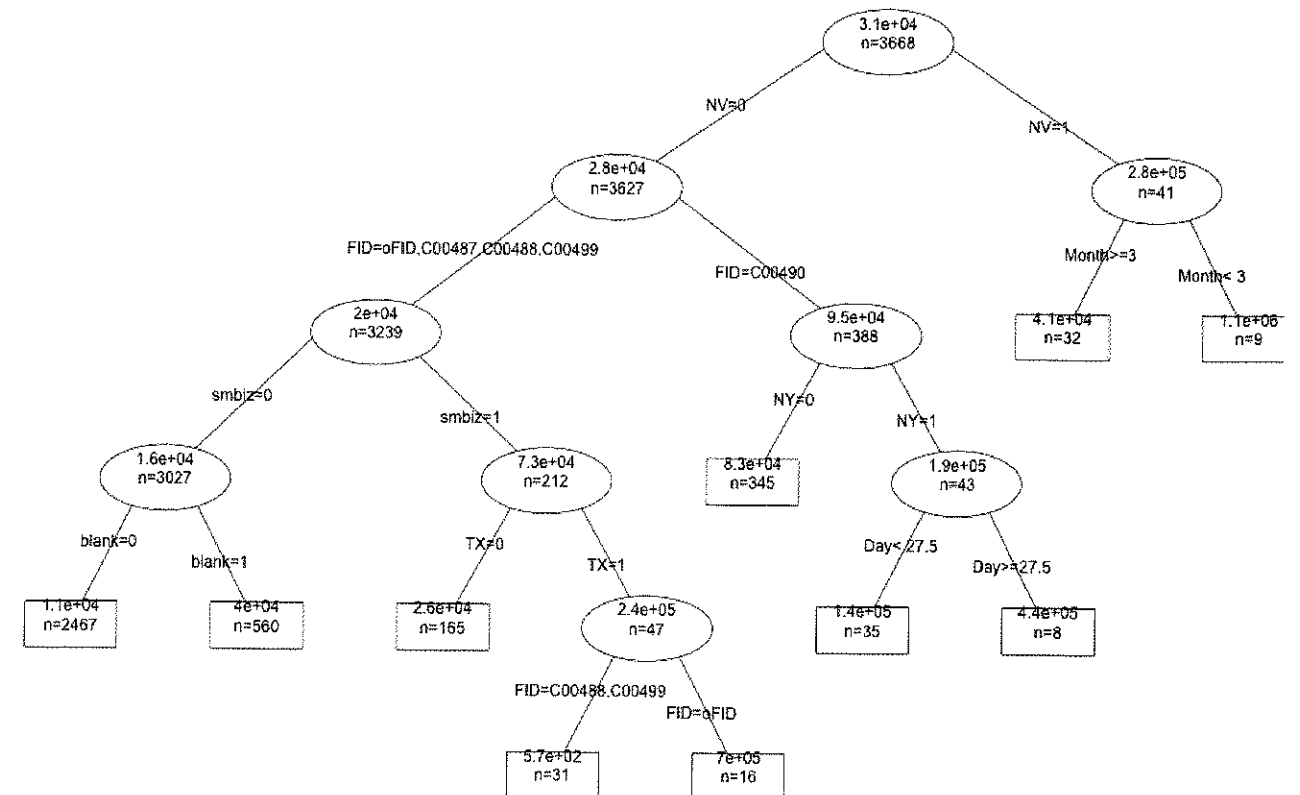
```

pdf
2

```


How do we run it? The graphical presentation.

Classifying SPAC Donation Size, 9 splits



What about exporting the results?

```
# will use a combination of list entries: frame, splits, and csplit
```

```
spac.tree9$frame[1:5, ]
```

##	var	n	wt	dev	yval	complexity	ncompete	nsurrogate
## 1	NV	3668	3668	1.438e+14	30936	0.037317	4	0
## 2	FID	3627	3627	9.400e+13	28138	0.016462	4	2
## 4	sembiz	3239	3239	8.088e+13	20113	0.016462	4	0
## 8	blank	3027	3027	2.897e+13	16381	0.002751	4	2
## 16	<leaf>	2467	2467	1.580e+13	10935	0.001581	0	0

井井井井

```
# frame is a matrix with 1 row per node of the tree
```

```
# row name corresponds to a unique node index
```

```
# var - name of the variable used in the split, or <leaf>
```

n - number of observations reaching the node

```
# yval - the fitted outcome value at the node
```

井井井井

```
spac.tree9$splits[1:5, ]
```

```
##          count ncat  improve  index adj
## NV          3668    2  0.017664    1.0  0
## FID         3668    5  0.012395    2.0  0
## Month       3668    1  0.005567    5.5  0
## sbmbiz      3668    2  0.004716    3.0  0
## retired     3668    2  0.003653    4.0  0
```

splits characterizes the splits making the regions R_m

```
# row name is the variable being split
```

count - the number of observations coming into the split

```
# ncat - number of categories of categorical variable, or 1 if the
# variable is numeric
```

```
# improve - the improvement in the objective using the split
```

```
# index - either the row number of the csplit matrix (for categorical
# variables), or the value of the optimal split (for numeric variables)
```

```
spac.tree9$csplit[1:5, ]
```

	[,1]	[,2]	[,3]	[,4]	[,5]
#					
# [1,]	1	3	2	2	2
# [2,]	1	3	1	1	1
# [3,]	1	3	2	2	2
# [4,]	3	1	2	2	2
# [5,]	1	3	1	1	1

has 1 row for each split on a categorical variable

the row number corresponds to index in `spac.treeell$split` above

each column is an ordered level of a categorical variable, up to the max levels of any categorical var

an entry of 1 - that level goes left in the split

3 - that level goes right in the split

2 - that level is not included in the split

What about exporting the results?

- To recreate a decision tree, you would at least extract the following columns of information:
 - `rownames(spac.tree9$splits)`
 - `spac.tree9$splits[, "count"]`, `spac.tree9$splits[, "index"]` and `spac.tree9$splits[, "ncat"]`
 - `spac.tree9$frame[, "var"]`, `spac.tree9$[, "n"]` and `spac.tree9$frame[, "yval"]`
 - `spac.tree9$csplit` corresponding to the rows given by "index" where "ncat" > 2 in "splits"
- The order of splits in "frame" are depth first, and left branch first
- Match between "frame" and "splits" by variable name and number of observations
 - since a variable can be split multiple times, and frame also includes competing and surrogate splits

Automatic Way to Select Tree Size

Can calculate contribution of split to decreasing objective $e(T)$ by

$$e_m = \frac{1}{N} \sum_{x_i \in R_m} (y_i - \bar{y}_m)^2$$

$$Imp_m = e_m - e_{ml} - e_{mr}$$

If $Imp_m \geq cp$ then accept the split, otherwise make m a terminal node

- $cp > 0$ is a tuning parameter, giving tree sizes as in "cptable"
- Actually a little trickier because the rule is applied in inverse order of depth

Solves the problem:

$$\min_T [e(T) + cp|T|]$$

- where $|T|$ is the number of terminal nodes of the tree

Automatic Way to Select Tree Size

- The entry "cptable" gives tree statistics for each cp
- "rel error" is the ratio of the objective, $e(T)$, to that of a single root tree
 - This is **always** decreasing with cp
- "xerror" is the average of 10 fold cross validation error
 - i.e. leave out 1/10th of the dataset,
 - train a size n tree on the other 9/10ths,
 - and compute $e(T)$ on the left out part
 - this is more useful for prediction, and not as useful to us for describing a dataset
 - can be thought of as a measure of **pervasiveness**
- Could consider a criteria that penalizes large trees
 - Not unreasonable: $N \times (relerror) + 2|T|$

Automatic Way to Select Tree Size

```
which.min(spac.tree$cptable[, 4])
```

```
1
1
```

gives a value of 1, meaning none of the splits are 'pervasize'

but using the criteria above, penalizing large trees

```
cpstat = dim(spac.data)[1] * spac.tree$cptable[, 3] + 2 * (spac.tree$cptable[, 2] + 1)
```

```
which.min(cpstat), , 3)
```

CP	nsplit	rel error	xerror	xstd
0.001	39.000	0.808	1.064	0.313

suggests a tree size with 39 splits

Advantages of Trees

1. **Fast** computations
2. **Invariant** under monotone transformations of variables
 - *Scaling doesn't matter!*
 - *Immune to outliers in x*
3. **Resistance** to irrelevant variables, so can throw lots of variables into it
4. **One tuning parameter** (tree size, or cp)
5. **Interpretable** model representation
6. **Handles missing data** by keeping track of surrogate, or highly correlated, backup splits at every node
7. Extends to **categorical outcomes** easily

Advantages of Trees

Accuracy

- $F(x)$ may not be piecewise constant (but decent overall approximation)
- Data Fragmentation (ok, if you have lots of data)
- $F(x)$ must involve high order interactions

Variance

- Each subsequent split depends on the previous ones, so an error in a higher split is propagated down.
- Small change in dataset can cause big change in tree
 - If you only have a random sample of a population, this can be a problem.
 - Not as much of an issue if you're describing a dataset

CART libraries outside of R: weka

- weka 3: Data mining software in JAVA
- <http://www.cs.waikato.ac.nz/ml/weka/>
- Relevant class `weka.classifiers.trees.J48`
- Simple command line syntax
 - `java weka.classifiers.trees.J48 -t data/weather.arff -i`
- ARFF is Attribute-Relation File Format and data format for weka
 - `weka.core.converters` package contains converters for usual data files
- Also call classes directly

```
import weka.core.Instances;
import weka.classifiers.Evaluation;
import weka.classifiers.trees.J48;
...
Instances train = ... // from somewhere
Instances test = ... // from somewhere //
train classifier Classifier cls = new J48();
cls.buildClassifier(train);
// evaluate classifier and print some statistics
Evaluation eval = new Evaluation(train);
eval.evaluateModel(cls, test);
System.out.println(eval.toSummaryString("", false));
```

- `weka.gui.treevisualizer.TreeVisualizer` class to visualize trees

RT libraries outside of R: orange

orange: Data mining through visual programming programming or Python scripting.

<http://orange.biolab.si/>

has proprietary tab-delimited data format

• Can import from csv, but is not very robust

• More info: [/Orange.data.formats/](#)

Relevant function: `Orange.regression.tree.TreeLearner(...)`

Visualizing trees: Orange renders trees in dot - plain text graph description language readable by both human and computer


```
tree.dot(file_name="0.dot", node_shape="ellipse",
leaf_shape="box")
```

CART libraries outside of R: opencv

- opencv: (Open Source Computer Vision) is a library of programming functions for real time computer vision, in C++
 - <http://opencv.willowgarage.com/wiki/>
 - Uses n-dimensional array class Mat to store and operate on data
 - [core_basic_structures.html#mat](#)
 - CvDTree class is an honest representation of CART algorithm
 - [ml_decision_trees.html](#)
 - *mushroom.cpp* example file demonstrates how to use decision trees

References

Elements of Statistical Learning. 2009. New York. Springer. xxii, 745 p. : ill. ; 24 cm.



Jerome Friedman's 315b course notes

Two solutions to Disadvantages (extra slides)

1. Boosted Trees, aka Forests, MART

- $F(x) = \sum_{k=1}^K a_k f(x; c_m^k, R_m^k)$
- Now each $f()$ is a tree, and $F()$ is a linear combination of trees
- Each tree can model an additive effect, or many low order interactions
- Variance of a combination of identically distributed objects is lower than any individual
- Disadvantage: loses decision tree interpretability unless K is small

2. Random Forests

- Similar to boosted trees, but now random subsets of the data are used for each tree
- Simpler to fit than boosted trees
- Accuracy is usually somewhere in between a single tree and boosted trees

How are Boosted Trees Interpreted? (extra slides)

Relative Importance

- $Imp_l^2 = Avg \left[\sum_{m=1}^M Imp_m I(var(m) = l) \right]$
- Average overall improvement of objective by variable l

Partial Dependence

- $pd(x_l) = E_{notl} [F(x_l, x_{notl})]$
- Predicted outcome using x_l , after averaging out the others